

# Accelerated Batch Rendering: A Software Implementation

Hongyu Teng (ht919)

27 June 2023

For better visual effect (including animated images), check the online version:

<https://docs.google.com/presentation/d/1ffdKxlpKJ1MEsorblpnoqY0HKOvmpQs-GNU2ygfQlqY/edit>

<https://github.com/JoeyTeng/jaxrender>

# Motivation

# Applications of RL (Reinforcement Learning)



Robots



Self-driving

SayCan: Grounding Language in Robotic Affordances. <https://say-can.github.io/>

AutoPilot, Tesla: [https://www.tesla.com/en\\_EU/autopilot](https://www.tesla.com/en_EU/autopilot)

# Applications in RL




Robots

SayCan: Grounding Language in Robotic Affordances. <https://say-can.github.io/>



Self-driving

AutoPilot, Tesla: [https://www.tesla.com/en\\_EU/autopilot](https://www.tesla.com/en_EU/autopilot)



Project 3D objects onto 2D sensor  
(not just usual camera, maybe IR, LiDAR)

# Applications in RL

- Use simulators to train agents
- Take **direct observation** from environment, or use internal states
  - Generate 2D observations using 3D env data?

# Applications in RL

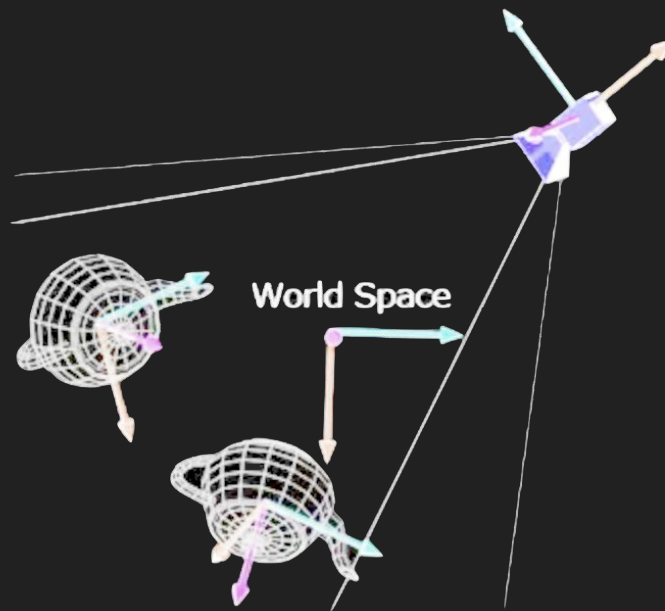
- Use simulators to train agents
- Take **direct observation** from environment, or use internal states
  - Generate 2D observations using 3D env data?

Use Renderers!

# Render



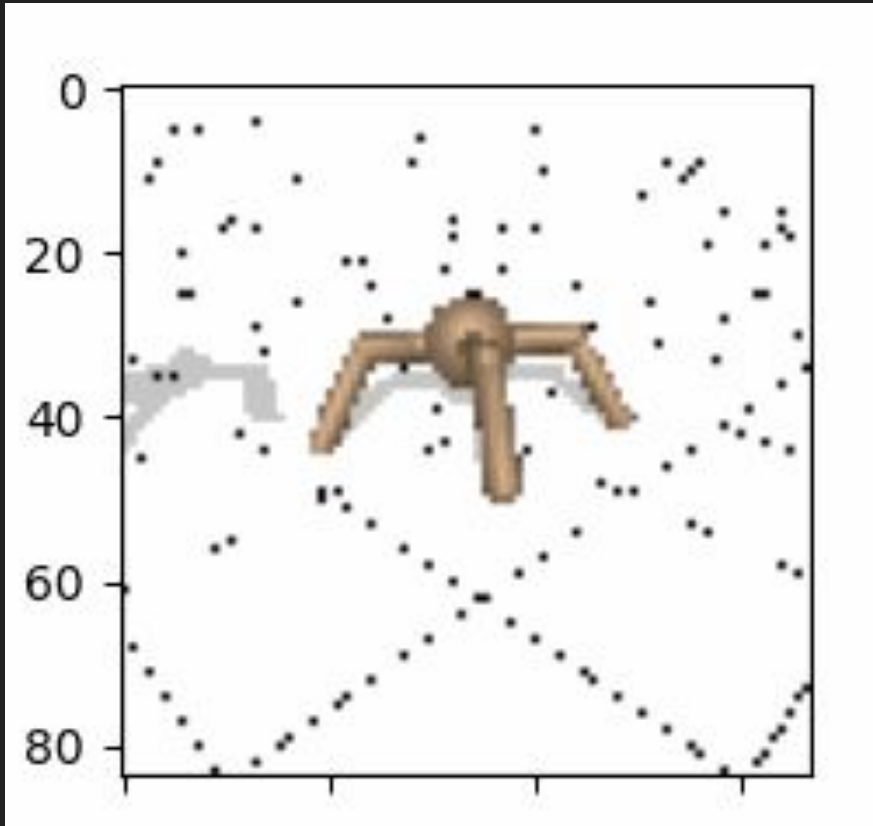
Project 3D objects onto 2D sensor,  
then discretely sample them



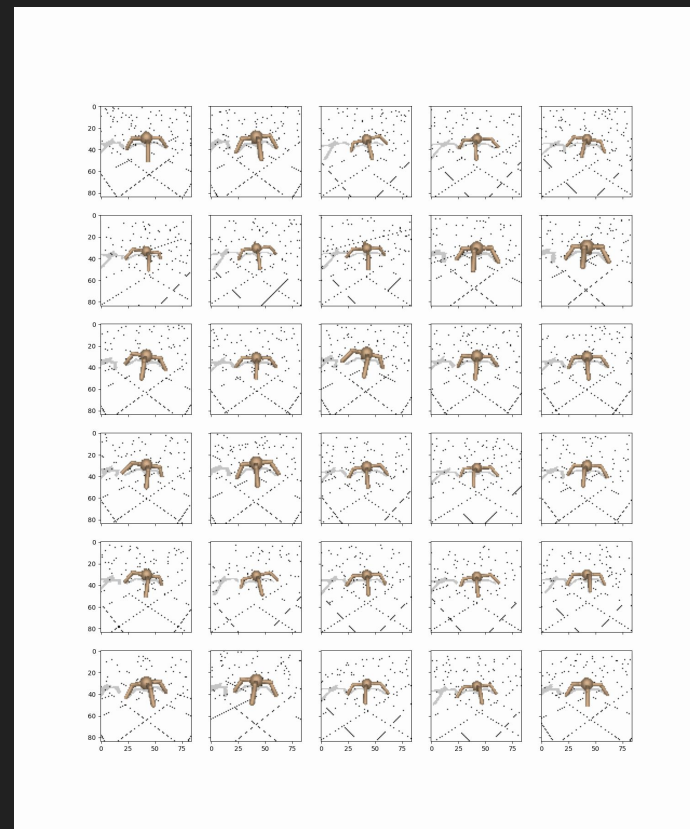

Similar to taking photos using  
camera in the 3D scene

# Exciting Examples





90 Frames Animation, 84 x 84, in **1.91 s (47fps)** (simulation + rendering)



30 Envs 90 Frames Animation, 84 x 84, in **5.26 s (513fps)** (simulation + rendering)

# Existing Solutions

- No batch rendering support, limit the throughput
- Not differentiable, limit the optimisation methods

# Existing Solutions vs Our Solution


- No batch rendering support, limit the throughput
- Not differentiable, limit the optimisation methods

Using JAX to implement a software renderer

- Native batch computation & parallel/distributed computation support
- Automatic differentiable
- Natively run on accelerators (GPU, TPU)

# Batch Rendering

Render several frames together to increase utilisation and throughput



X  
P  
U




Batch Rendering (6 per batch)


Sequential Rendering (1-by-1)

# Batch Rendering

Render several frames together to increase utilisation and throughput



Batch Rendering (6 per batch)




Sequential Rendering (1-by-1)


# Batch Rendering

Render several frames together to increase utilisation and throughput

Compiler automatically divide big batch into smaller at op-level, and schedule them. This is done in compile time. (OOM still possible)



Capability 6 per batch




Capability 3 per batch

# Batch Rendering

Render several frames together to increase utilisation and throughput

Compiler automatically divide big batch into smaller at op-level, and schedule them. This is done in compile time. (OOM still possible)

Easy distributed batch rendering on homogeneous XPU's at frame level



Render 12 frames in 2 XPU's in parallel, 6 frames each



# JAX: Autograd and XLA

The Numerical Computation Framework Used

# JAX: Advantages and Limitations

- |   |  |
|---|--|
| ✓ Accelerator support (GPU, TPU)  | ✗ Heavy dispatch overhead (asynchronised)  |
| ✓ Simple JIT/AOT <sup>[1]</sup> compilation   | ✗ Long compilation time <ul style="list-style-type: none"><li>○ Quadratic with the number of ops<sup>[3]</sup></li><li>○ Must compile whole function at once</li><li>○ Change input shape causes re-compilation</li></ul>                                  |
| ✓ Simple automatic vectorisation  |  |
| ✓ Simple parallelisation & distribution   |  |
| ✓ Numpy-like API + some Sci-Py functions  | ✗ Statically shaped inputs & intermediate values   |
| ✓ Autograd: automatic differentiation   | ✗ Auto-parallelisation/distribution requires homogeneous device and data (shape)   |
| ✓ Type-hinting support <sup>[2]</sup>   |  |
| ~ Functional programming style: <ul style="list-style-type: none"><li>○ Composable function transformations</li><li>○ Pure functions + immutable data</li></ul> | ✗ Limited data-dependent control flow <ul style="list-style-type: none"><li>○ data must be static (known at compile time)</li><li>○ or, requires synchronisation with host</li><li>○ or, all branches are executed unconditionally<sup>[4]</sup></li></ul> |

[1]: JIT: Just-in-Time; AOT: Ahead-of-Time

[2]: Typing support through community solution jaxtyping: <https://github.com/google/jaxtyping>

[3]: Roughly estimation, see <https://github.com/google/jax/discussions/3478>

[4]: Under transformation. <https://github.com/google/jax/pull/16335>

# JAX: Advantages and Limitations (highlighted main ones)

- ✓ Accelerator support (GPU, TPU)
- ✓ Simple JIT/AOT<sup>[1]</sup> compilation
- ✓ Simple automatic vectorisation
- ✓ Simple parallelisation & distribution
- ✓ Numpy-like API + some Sci-Py functions
- ✓ Autograd: automatic differentiation
- ✓ Type-hinting support<sup>[2]</sup>
- ~ Functional programming style:
  - Composable function transformations
  - Pure functions + immutable data
- ✗ Heavy dispatch overhead (asynchronised)
- ✗ Long compilation time
  - Quadratic with the number of ops<sup>[3]</sup>
  - Must compile whole function at once
  - Change input shape causes re-compilation
- ✗ Statically shaped inputs & intermediate values
- ✗ Auto-parallelisation/distribution requires homogeneous device and data (shape)
- ✗ Limited data-dependent control flow
  - data must be static (known at compile time)
  - or, requires synchronisation with host
  - or, all branches are executed unconditionally<sup>[4]</sup>

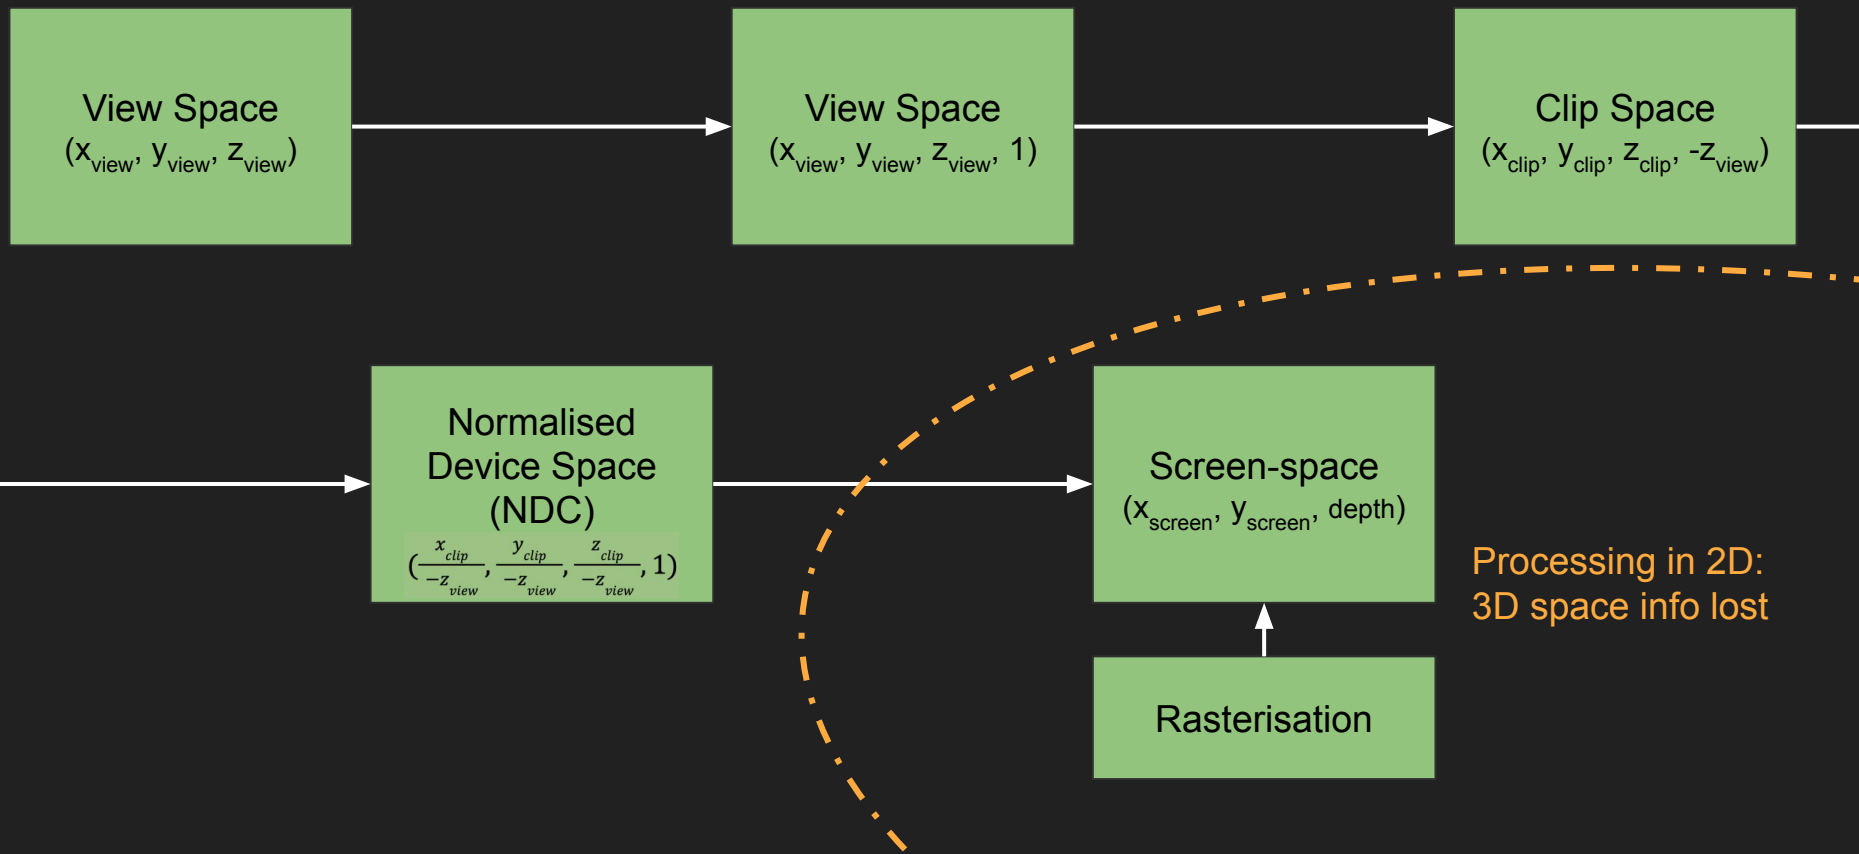
[1]: JIT: Just-in-Time; AOT: Ahead-of-Time

[2]: Typing support through community solution jaxtyping: <https://github.com/google/jaxtyping>

[3]: Roughly estimation, see <https://github.com/google/jax/discussions/3478>

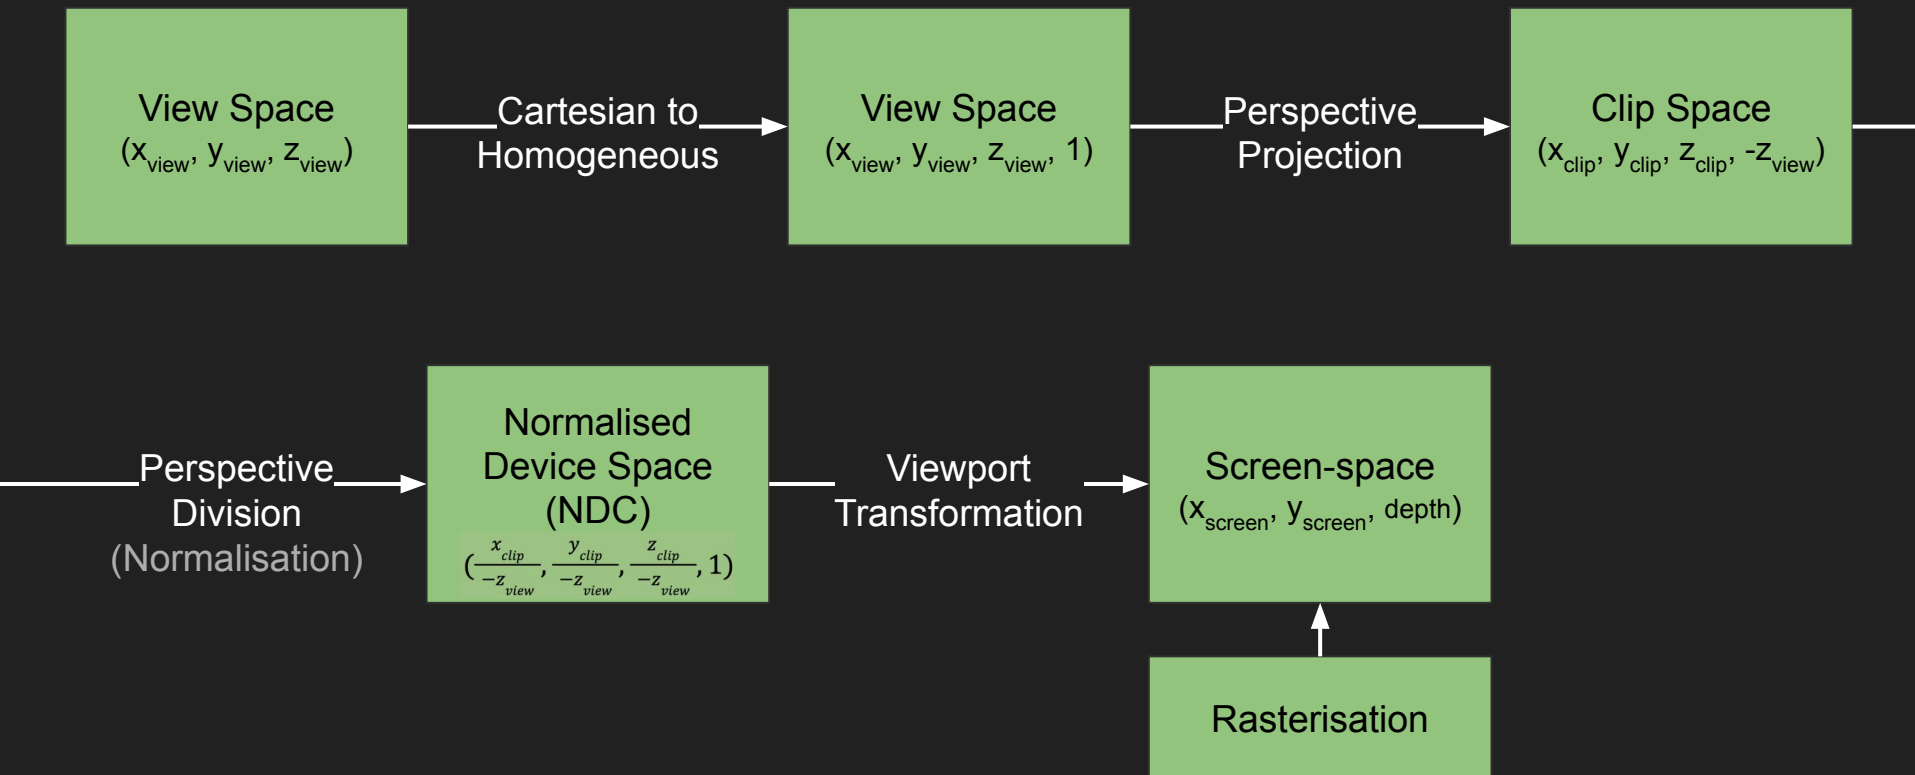
[4]: Under transformation. <https://github.com/google/jax/pull/16335>

# From 3D Models to Screen Pixels



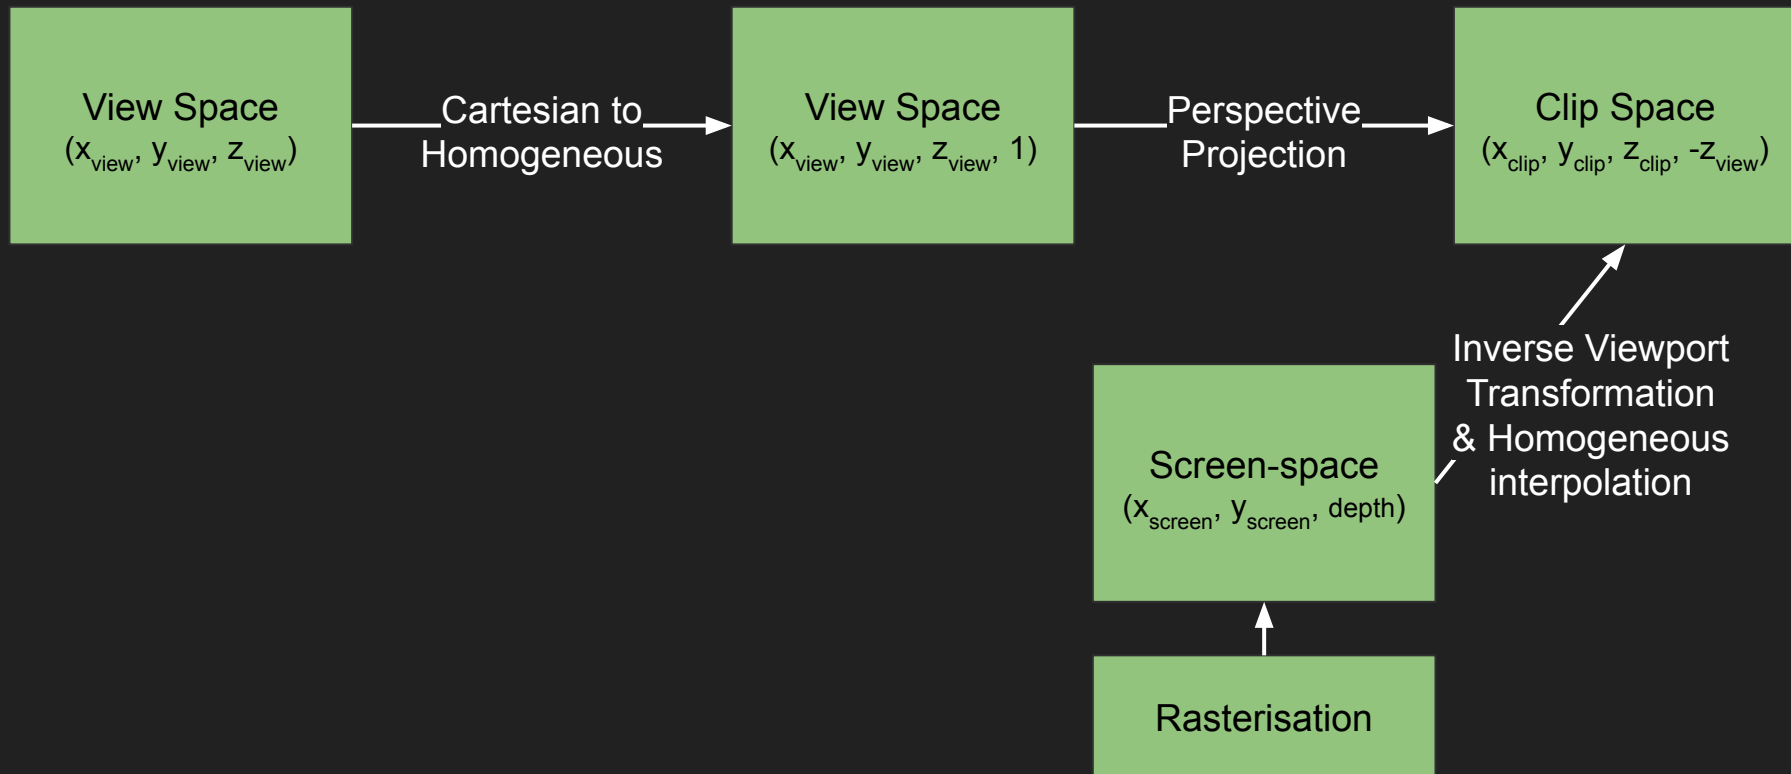
## OpenGL Transformation of coordinates from View space to screen space

Typically the coordinates will be transformed from world space to view space via a *view matrix*, but that is defined by user in *vertex shader*, and is not part of OpenGL specification.



## OpenGL Transformation of coordinates from View space to screen space

Typically the coordinates will be transformed from world space to view space via a *view matrix*, but that is defined by user in *vertex shader*, and is not part of OpenGL specification.

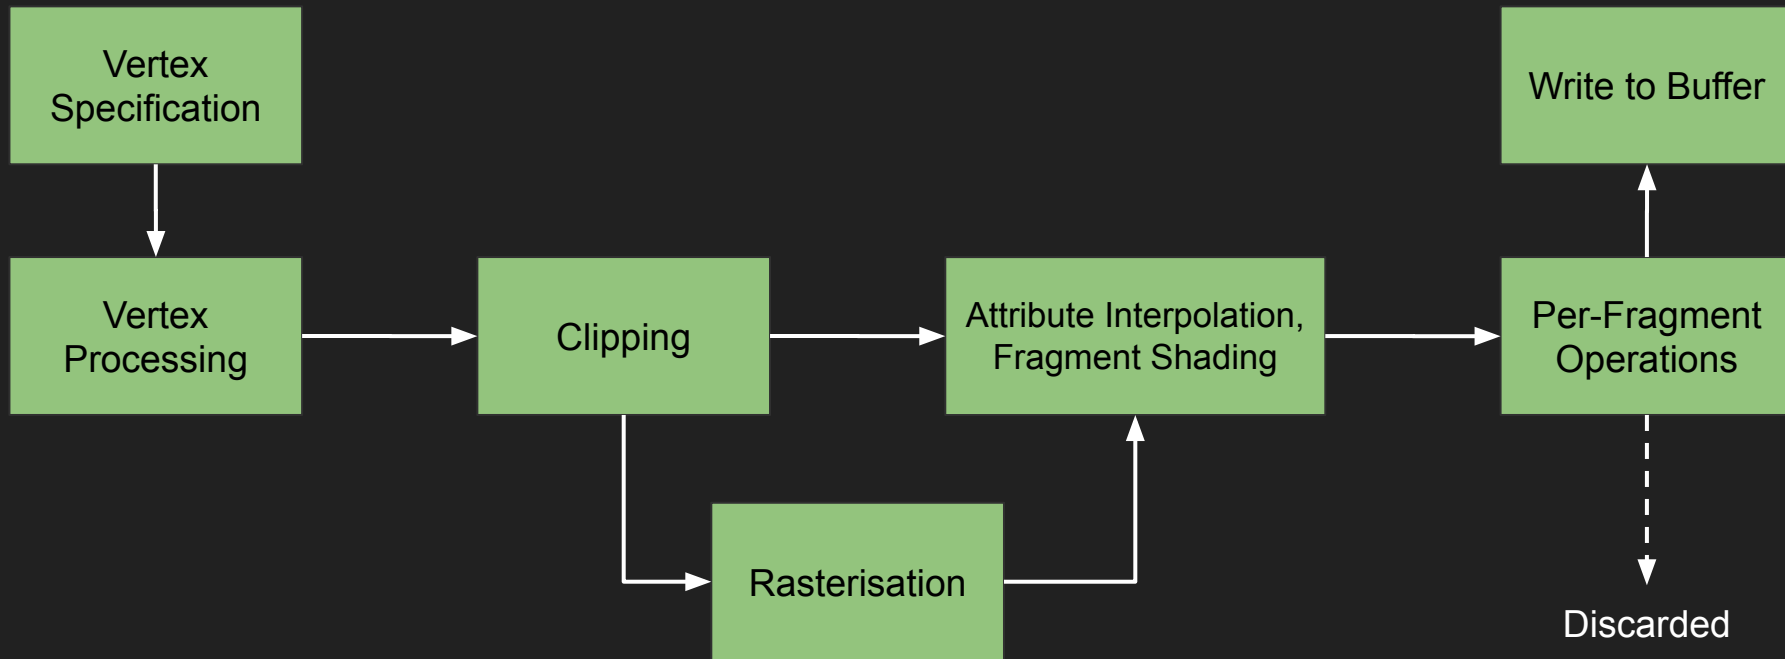


## Our Transformation of coordinates

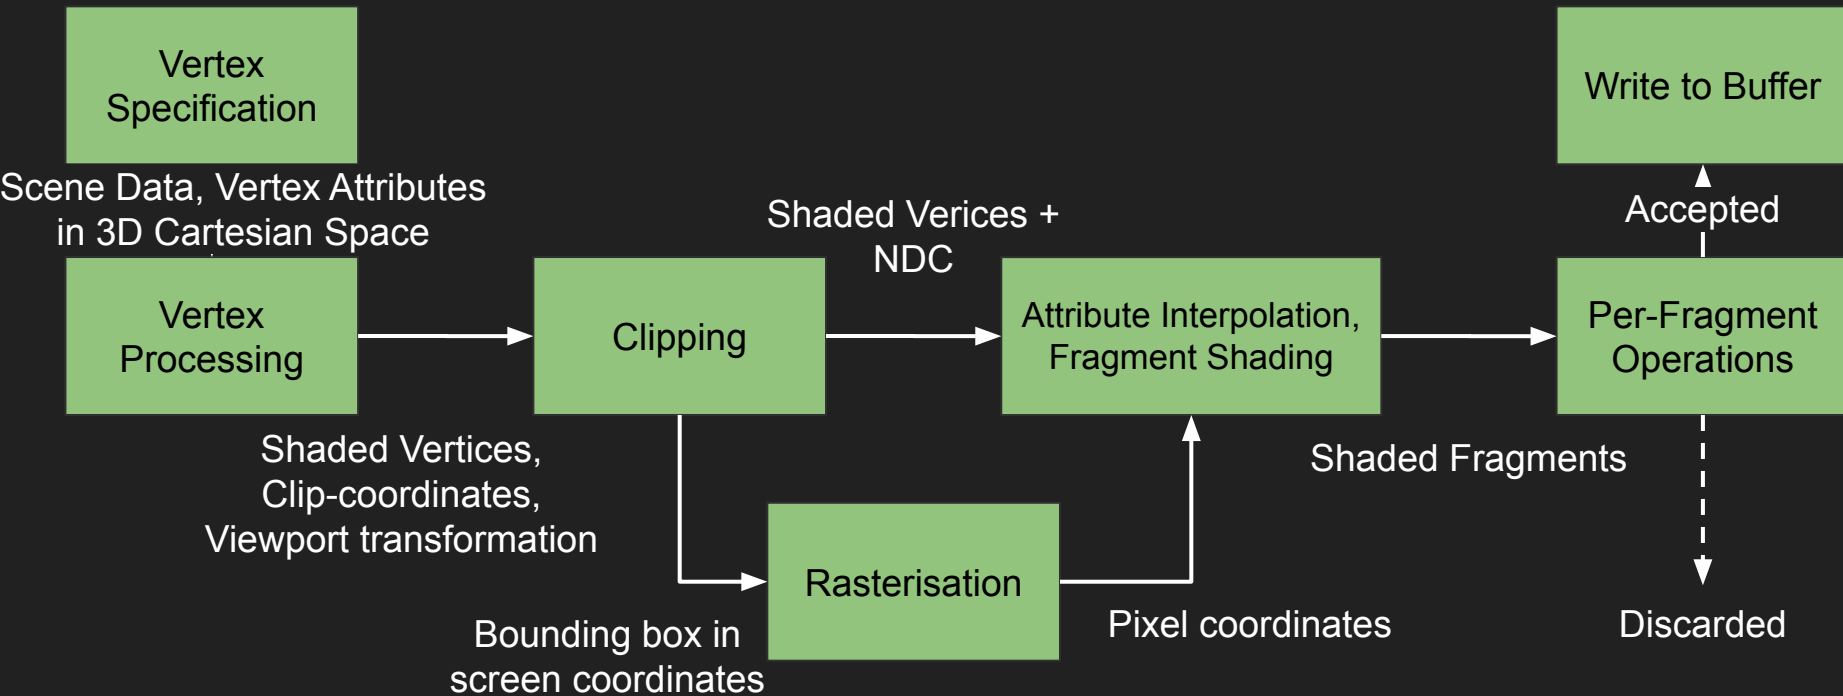
Convert pixels into 3D space to interpolate

# Graphics Pipeline (OpenGL)

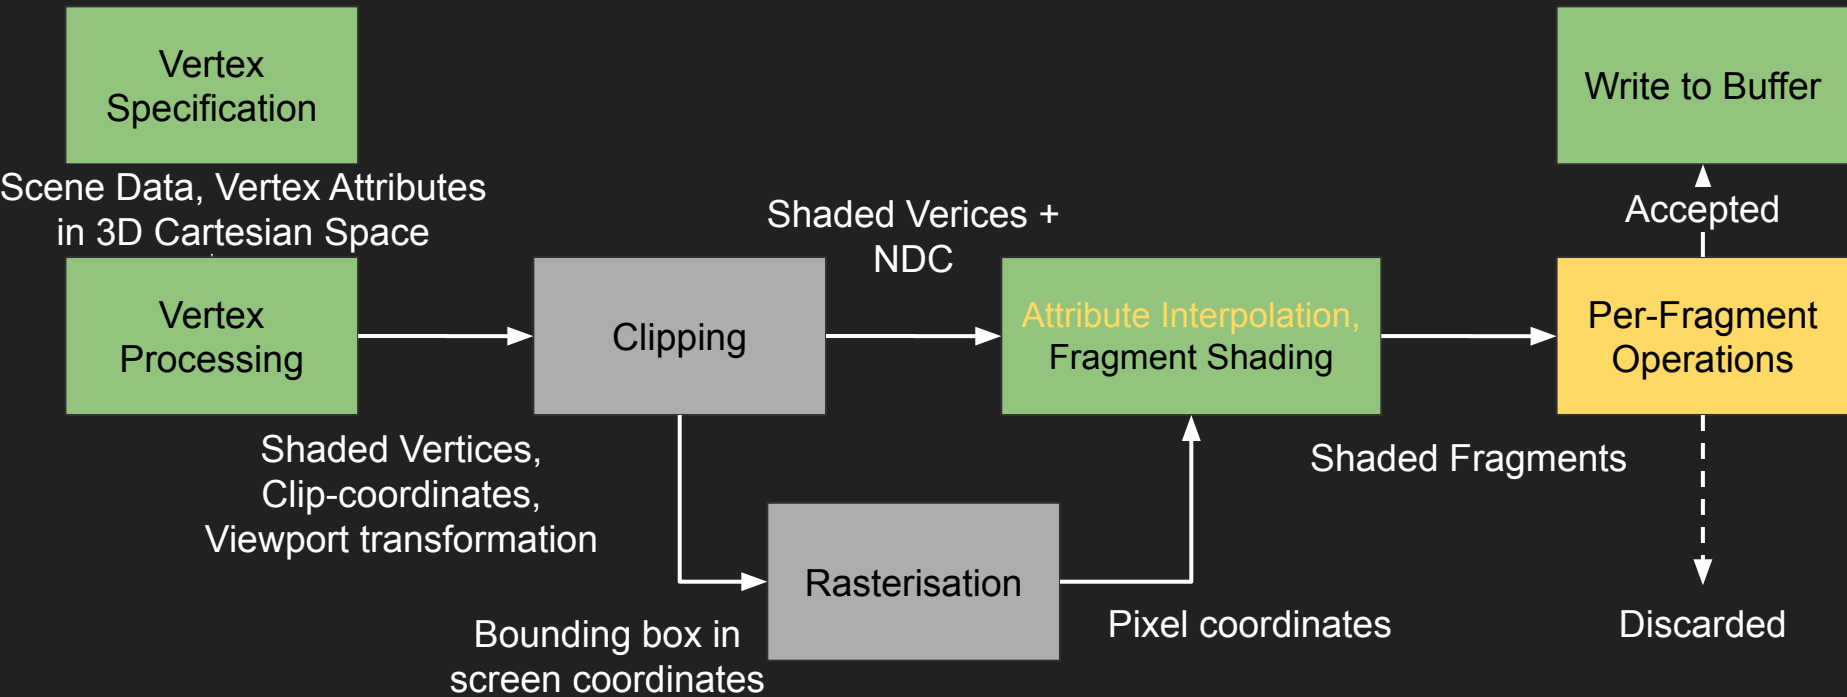




## Overview of OpenGL Pipeline (Compulsory Stages)

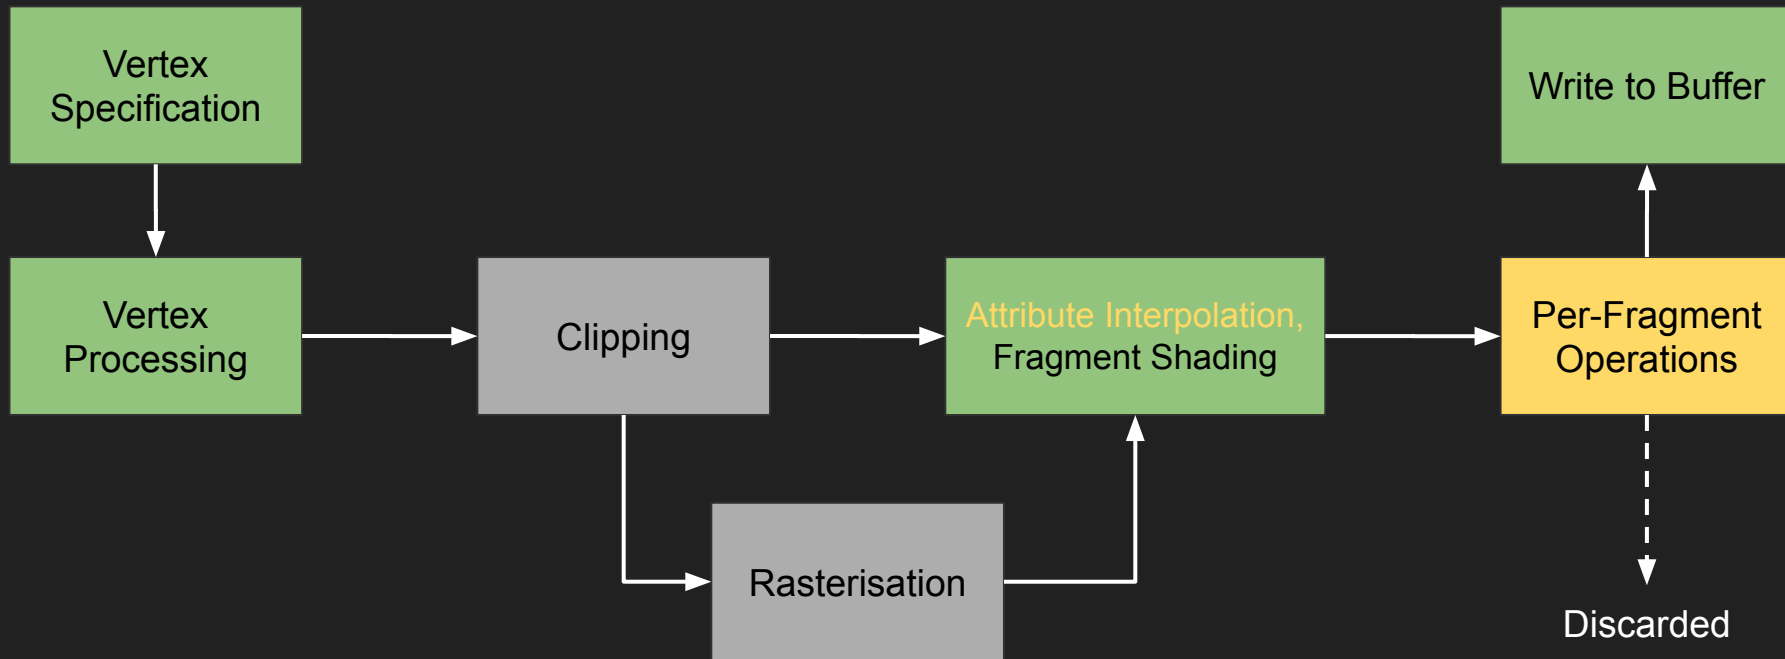


## Overview of OpenGL Pipeline with Data Flow



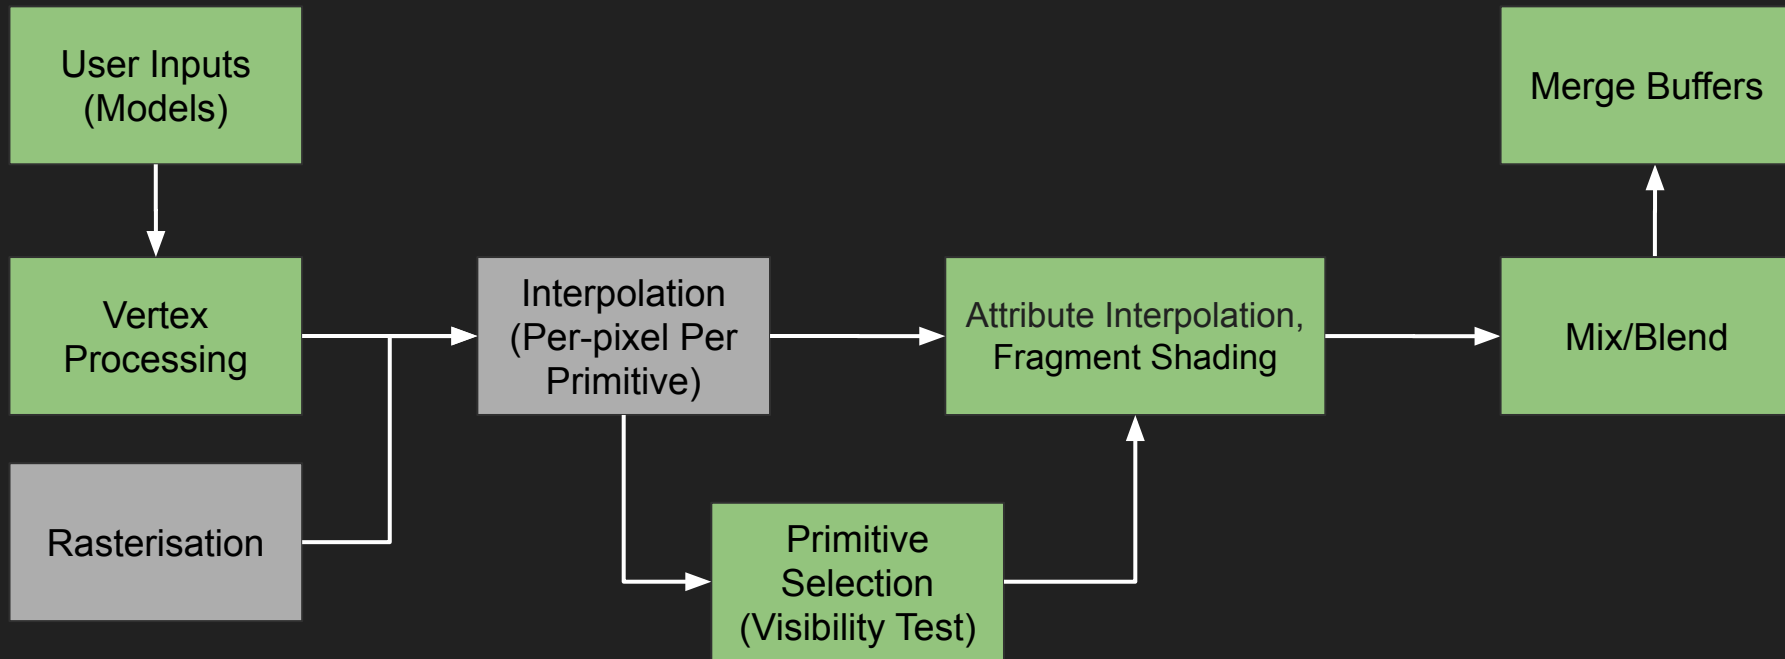
## Overview of OpenGL Pipeline with Data Flow

User-defined  
Customisable by choosing  
between given options  
Fixed, not customisable



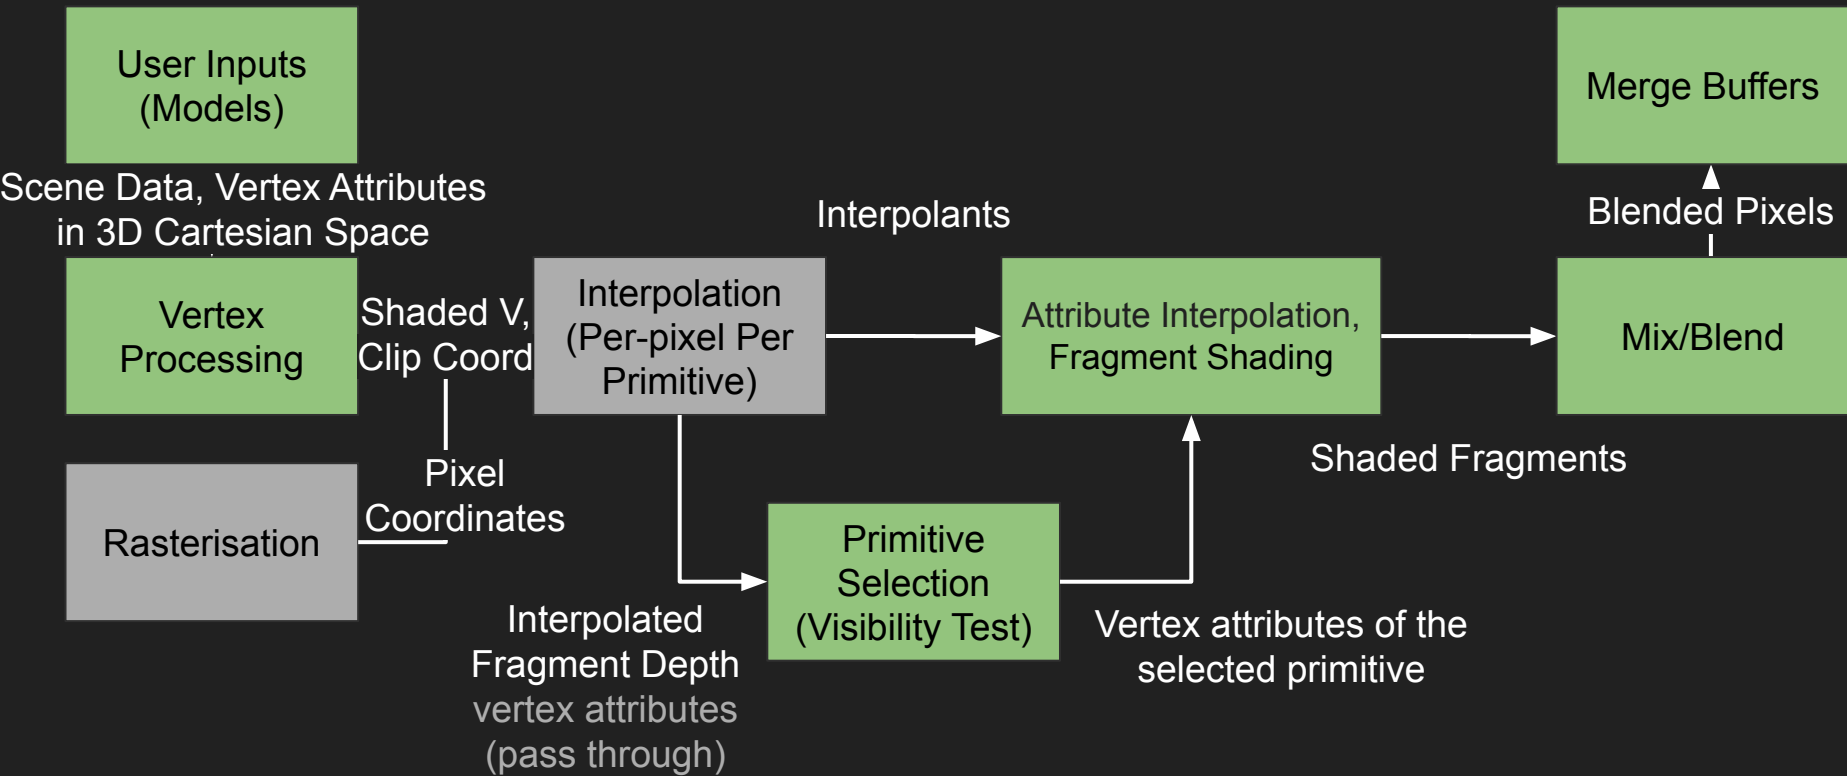
## Overview of OpenGL Pipeline (Compulsory Stages)

User-defined  
Customisable by choosing  
between given options  
Fixed, not customisable



## Overview of JaxRenderer Pipeline

User-defined with default given  
Fixed, not customisable



## Overview of JaxRenderer Pipeline

User-defined with default given  
Fixed, not customisable

# Changes Made

## Ours

1. Use homogeneous interpolation + customisable interpolation shader

## OpenGL

1. Perspective-corrected screen-space interpolation & clip stage + interpolation keywords

# Changes Made

## Ours

1. Use homogeneous interpolation + customisable interpolation shader

## OpenGL


1. Perspective-corrected screen-space interpolation & clip stage + interpolation keywords

*Why can we skip clipping stage?*




# Clipping

- **Necessary** for the correctness of perspective division (normalise homogeneous coordinates)
- Perspective projection:
- World space  $\Rightarrow$  Eye (clip) space
- The homogeneous coordinates must be normalised to extract the equivalent 3D cartesian coordinates




Eye Space



NDC

without clipping



NDC

with clipping

Triangles with vertex behind camera will have the sign of  $z$  flipped during perspective division (without clipping). Clipping solves this issue

Reference: <https://zhuanlan.zhihu.com/p/102758967>


# Clipping

- **Necessary** for the correctness of perspective division (normalise homogeneous coordinates)
- Perspective projection:
- World space  $\Rightarrow$  Eye (clip) space
- The homogeneous coordinates must be normalised to extract the equivalent 3D cartesian coordinates


$$(x_{view}, y_{view}, z_{view}, 1)$$

$$\rightarrow (x_{clip}, y_{clip}, z_{clip}, -z_{view})$$

$$\rightarrow \left( \frac{x_{clip}}{-z_{view}}, \frac{y_{clip}}{-z_{view}}, \frac{z_{clip}}{-z_{view}}, 1 \right)$$




Eye Space



NDC

without clipping



NDC


with clipping

Triangles with vertex behind camera will have the sign of  $z$  flipped during perspective division (without clipping). Clipping solves this issue

Reference: <https://zhuanlan.zhihu.com/p/102758967>

# Homogeneous Interpolation

- Perspective-correct interpolation (interpolating in clip-space)
- Correctly handles all homogeneous coordinates without normalisation
- No clipping stage needed



# Changes Made: Summary

## Ours

1. Use homogeneous interpolation + customisable interpolation shader
2. Add primitive selection stage to test visibility earlier, across all triangles for each pixel
3. Rasterise on whole canvas for all selected triangles per pixel altogether (static shape)
4. Customisable mixing/blending stage

## OpenGL

1. Perspective-corrected screen-space interpolation & clip stage + interpolation keywords
2. Only do z buffering at the end, and has earlier triangles will not be rejected if further away than later shaded triangles
3. Rasterise on each triangle within their bounding box in tiles
4. Carefully design triangle rendering order and blending equation for desired mixing effect

# Quick Demo

1. Simple API
2. Customised Shaders
3. Differentiability

# Conclusion

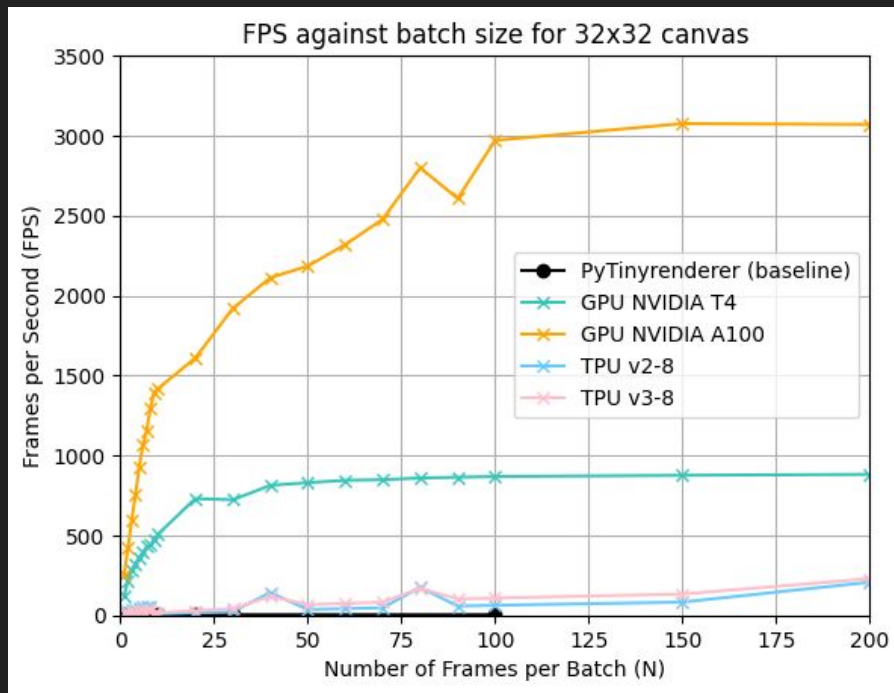
## Ours

1. Use homogeneous interpolation + customisable interpolation shader
  2. Add primitive selection stage to test visibility earlier, across all triangles for each pixel
  3. Rasterise on whole canvas for all selected triangles per pixel altogether
  4. Customisable mixing/blending stage
- 
1. Batch Rendering (higher throughput)
  2. More flexible, more convenient API
  3. Differentiable Rendering
  4. XPU support (TPU, GPU (CUDA, ROCm), Apple Silicon)

## OpenGL


1. Perspective-corrected screen-space interpolation & clip stage + interpolation keywords
2. Only do z buffering at the end, and has earlier triangles will not be rejected if further away than later shaded triangles
3. Rasterise on each triangle within their bounding box in tiles
4. Carefully design triangle rendering order and blending equation for desired mixing effect

# Performance: Batch Size (Ant, 3276 Triangles)



Far exceeds baseline (CPU rendering)  
throughput growth as batch size grow

# Performance: Canvas Size (Ant, 3276 Triangles)



6.8x more pixels, 1/3 throughput

Large batch size is more meaningful with smaller canvas





Phong Reflection Model + Shadow, 30 frames 1920 x 1080, 2492 triangles, in 9.25s.

Thank you!

# Backup Slides

# JAX: Vectorisation vs Parallelisation

- Ⓥ Lowering the vectorisation into each lowest op (operator), then compile
- Ⓥ Execute in single device
- Ⓥ Do not support input with different shapes along batch axis
- Ⓥ Input batch size will be lowered to the op's batch size, may lead to unexpected high memory usage for intermediate values
- Ⓟ Replicate the compiled program to all devices, then execute them together, act in SPMD (Single-Program-Multiple-Data) pattern.
- Ⓟ Execute in multiple devices in parallel
- Ⓟ Input must be of same shape across devices
- Ⓟ Input batch axis size must be a multiple of the number of devices
- Ⓟ Devices must be homogeneous